



OnboardAI

AI Employee Pre-boarding & Onboarding Command Center

AI For All 2.0 • Cowork Capstone • Use Case 2 — People Ops & Onboarding Concierge

Amol Parbat • 25 Sep 2026

The problem

A new hire's first two weeks are scattered across people, tools and documents

Repetitive questions

People Ops answers the same leave, hours and document questions again and again — by hand.

Generic answers

HR bots answer confidently but not from THIS company's policies — so answers can be wrong.

Blockers found late

A BGV discrepancy or missing bank form is noticed only when payroll slips.

Result: slow, inconsistent onboarding and a poor first impression for the new hire.

The solution: one AI command center

Grounded in company data — every answer comes from a tool, not LLM memory

Welcome

Personal note (Groq) +
Gmail email + Calendar
invite for virtual roles

Answer

Policy answers from
Aarohan's own policies,
with source & confidence

Track

Progress %, blockers and
next priorities per new hire


Protect

Pay / legal / grievance
questions escalated to a
human

Users: People Ops coordinators (primary) • Hiring managers • New hires (via assistant)


OnboardAI overview

Four tabs — simple enough for a non-technical HR user



- Dashboard — new hires, average progress, pending tasks
- New Hire Welcome — generate, review, send
- Policy Assistant — Quick Lookup, Policy Intelligence, AI chat
- Onboarding Progress — status, tasks, blockers
- Colours taken from the Aarohan logo; works on mobile

Simplest thing that works: static page + one serverless function



Gmail · Google Calendar · Google Drive · Gemini embeddings (inside n8n)

The LLM writes and decides — tools supply the facts

Personalised writing

Groq drafts a warm, human welcome note using the hire's role, department, start date and work mode.

Grounded Q&A (RAG)

Semantic search over 38 policy chunks; confidence gating decides; Groq only rephrases the evidence.

Agentic tool use

Chat uses Groq tool-calling: it picks `get_policy`, `answer_policy_question`, `get_onboarding_status` or `list_new_hires`.

Guardrails

Pay, legal, grievance, termination → escalated to a human before the LLM runs. No cross-employee data sharing.

4 custom MCP servers — used by the web app AND by Claude Cowork

MCP endpoint (n8n)	Tool	What it returns
onboardai-policy-lookup	get_policy	Policy text + Google Drive source link
onboardai-policy-intelligence	answer_policy_question	Answer, source section, similarity, confidence, human-review flag
onboardai-onboarding-progress	get_onboarding_status	% complete, status, blockers, next priorities
onboardai-new-hire-welcome	onboard_new_hire	Email sent / calendar created (virtual only)

The app speaks MCP directly: initialize → tools/call over Streamable HTTP (JSON or SSE). Every result in the UI shows which MCP tool answered.

Each tool = thin MCP wrapper + a Logic workflow that does the real work

- MCP Server Trigger → "Call n8n Workflow Tool" → Engine (Logic) sub-workflow
- Engines read n8n Data Tables (Employees, Tasks, Documents, Policies)
- Welcome engine: Gmail send → IF virtual → Google Calendar event
- Intelligence engine: Gemini embedding → Supabase pgvector top-5 → deterministic evidence & confidence rules
- All 9 workflows + data exported to n8n-backup/ (credentials stripped) before the trial ends

Why this matters

Business rules (what counts as "Blocked", when to ask a human) live in visible, auditable workflows — not hidden inside a prompt.

Persistent data + vector search

onboarding_employees / tasks

New-hire list and task checklist. New hires added in the app get a default checklist automatically.

policy_chunks (pgvector)

38 chunks from 8 Google Drive policy docs, 768-dim embeddings, cosine search via the Intelligence MCP.

onboarding_activity


Audit log: welcome sent (email + calendar status), chat tool usage.

Row-Level Security on every table • accessed only from the server function


The only LLM provider — fast, free tier, key never leaves the server

- Model: llama-3.3-70b-versatile (configurable via GROQ_MODEL)
- Welcome drafting · evidence → plain-English answers · tool-calling chat agent (max 4 tool rounds)
- GROQ_API_KEY stored as a Netlify secret env var; browser only calls /api/*
- No fallback providers — one predictable, low-cost dependency


Screens from the local test run (same UI as deployed)



Welcome: draft → send



Policy answers + chat with tool trace



Progress: status, tasks, blockers

11 Live demo flow (5 minutes)

- 1 Dashboard: 5 new hires, average progress, pending tasks
- 2 Welcome Devika (virtual): Groq note → Send → Gmail email + Calendar invite
- 3 Ask "What is the leave policy?" → Quick Lookup with Drive link
- 4 Ask "Day-1 schedule for remote hires" → Intelligence: high confidence + source
- 5 Progress for Arjun (EMP-1003): Blocked — BGV discrepancy
- 6 Chat "Can I get a salary hike?" → escalated to a human

Expected impact — to be measured in a pilot, not yet proven

Time back for People Ops

Routine policy questions and status checks answered self-serve; welcome + invite in one click.

Consistent, correct answers

Only from company policy, with source; low-confidence answers flagged for a human.

Earlier risk detection

BGV, document and overdue-task blockers surfaced before Day 1 — protects first payroll.

Pilot metrics to track: % questions answered without HR · time from offer to "ready for Day 1" · # hires with payroll delayed

- API keys only in Netlify environment variables — never in the browser or GitHub (.env.example has names only; automated secret scan)
- n8n backups exported with all credentials stripped
- Input validation (employee ID format, question length) and HTML-escaped output
- Row-Level Security enabled on all Supabase tables
- Human escalation for pay, legal, grievance and termination topics
- Fictional data only — demo RLS and open MCP endpoints must be tightened before real use

Honest limitations

What it does not do yet

- n8n cloud trial ends in about 8 days — workflows backed up, migration planned
- No login or roles yet: anyone with the link can view demo data and send test welcomes
- Employee data split between n8n Data Tables and Supabase
- Policy Intelligence embeddings use Gemini inside n8n (stored vectors depend on it)
- Welcome email body is a fixed n8n template; the Groq note is shown in the app
- Limitation I would fix next: unify data + add Supabase Auth

15 Future roadmap

Week 1

Move the 4 MCP tools to Supabase Edge Functions; one data store

Week 2

Supabase Auth with roles: new hire / manager / People Ops

Week 3

Day-3 proactive nudges (Slack/email) + weekly manager digest

Later

Document upload & verification, Slack buddy assignment, time-to-productive analytics